

# Propositions a decentralized IaaS extending OpenStack

Adrien Lebre, Jonathan Pastor  
ASCOLA Research Group  
Mines Nantes / Inria / LINA UMR 6241  
Nantes, France

Frédéric Desprez  
Corse Research Group  
Inria / LIG UMR 5217  
Grenoble, France

## ABSTRACT

Instead of the current trend consisting of building larger and larger data centers (DCs) in few strategic locations, the DISCOVERY initiative<sup>1</sup> proposes to leverage any network point of presences (PoP, *i.e.*, a small or medium-sized network center) available through the Internet. The key idea is to demonstrate a widely distributed Cloud platform that can better match the geographical dispersal of users. This involves radical changes in the way resources are managed, but leveraging computing resources around the end-users will enable to deliver a new generation of highly efficient and sustainable Utility Computing (UC) platforms, thus providing a strong alternative to the actual Cloud model based on mega DCs (*i.e.* DCs composed of tens of thousands resources).

Critical to the emergence of such distributed Cloud platforms is the availability of appropriate operating mechanisms. Although, some of protagonists of Cloud federations would argue that it might be possible to federate a significant number of micro-Clouds hosted on each PoP, we emphasize that federated approaches aim at delivering a brokering service in charge of interacting with several Cloud management systems, each of them being already deployed and operated independently by at least one administrator. In other words, current federated approaches do not target to operate, remotely, a significant amount of UC resources geographically distributed but only to use them. The main objective of DISCOVERY is to design, implement, demonstrate and promote a unified system in charge of turning a complex, extremely large-scale and widely distributed infrastructure into a collection of abstracted computing resources which is efficient, reliable, secure and friendly to operate and use.

After presenting the DISCOVERY vision, we explain the different choices we made, in particular the choice of revising the OpenStack solution leveraging P2P mechanisms. We believe that such a strategy is promising considering the architecture complexity of such systems and the velocity of open-source initiatives.

## Keywords

Cloud computing, IaaS Architecture, OpenStack, NoSQL, Peer to Peer

## 1. INTRODUCTION

<sup>1</sup><http://beyondtheclouds.github.io>

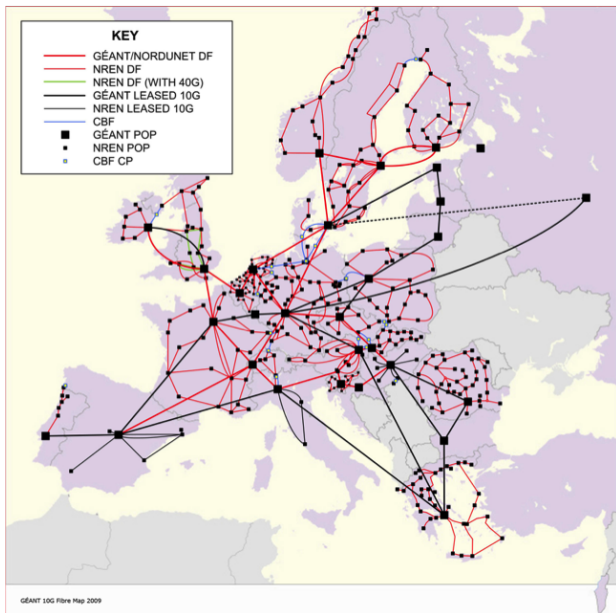
To satisfy the escalating demand for Cloud Computing (CC) resources while realizing economy of scale, the production of computing resources is concentrated in mega data centers (DCs) of ever-increasing size, where the number of physical resources that one DC can host is limited by the capacity of its energy supply and its cooling system. To meet these critical needs in terms of energy supply and cooling, the current trend is toward building DCs in regions with abundant and affordable electricity supplies or in regions close to the polar circle to leverage free cooling techniques [11].

However, concentrating Mega-DCs in only few attractive places implies different issues. First, a disaster<sup>2</sup> in these areas would be dramatic for IT services the DCs host as the connectivity to CC resources would not be guaranteed. Second, in addition to jurisdiction concerns, hosting computing resources in a few locations leads to useless network overheads to reach each DC. Such overheads can prevent the adoption of the UC paradigm by several kind of applications such as mobile computing or big data ones.

The concept of micro/nano DCs at the edge of the backbone [12] is a promising solution to address the aforementioned concerns. However, operating multiple small DCs breaks somehow the idea of mutualization in terms of physical resources and administration simplicity, making this approach questionable. One way to enhance mutualization is to leverage existing network centers, starting from the core nodes of the backbone to the different network access points (*a.k.a.* PoPs – Points of Presence) in charge of interconnecting public and private institutions. By hosting micro/nano DCs in PoPs, it becomes possible to mutualize resources that are mandatory to operate network/data centers while delivering widely distributed CC platforms better suited to cope with disasters and to match the geographical dispersal of users and their needs (see Figure 1).

A preliminary study has established the fundamentals of such an *in-network distributed cloud* referred by the consortium as the *Locality-Based Utility Computing* (LUC) concept [2]. However, the question of how operating such an infrastructure is still under investigations. Indeed, at this level of distribution, latency and fault tolerance become primary concerns, and collaboration between servers of differ-

<sup>2</sup>On March 2014, a large crack has been found in the Wapum Dam leading to emergency procedures. This hydrolic plan supports the utility power supply to major data centers in central Washington.



**Figure 1: The European GÉANT backbone**

GÉANT is the federation of all European Research and Educational Networks. Each black square corresponds to one network point of presence (*a.k.a.* a PoP) that can host a nano/micro DC.

ents location must be organized wisely.

In this vision paper, we discuss some key-elements that motivate our choices to design and implement the *LUC Operating System*, a system in charge of turning a LUC infrastructure into a collection of abstracted computing facilities that are as convenient to administrate and use as available Infrastructure-as-a-Service (IaaS) managers [7, 21, 22]. We explain, in particular, why federated approaches [3] are not satisfactory and why designing a fully distributed system that operates all resources makes sense.

As the capabilities of the LUC OS are similar to existing IaaS managers and because it would be a non-sense technically speaking to develop the LUC OS from scratch we chose to revise the OpenStack solution [22], leveraging P2P mechanisms. Our current efforts focus on the validation of a distributed version of the *Nova* service on top of Grid’5000 [1]. Historically, *Nova* relies on a MySQL centralized database, preventing it to natively scale beyond one site. To reach such a goal, we replaced the MySQL component by the REDIS backend, a distributed key/value store. Such a modification enables us to deploy several *Nova* controllers on distinct sites giving the illusion that there was only one global infrastructure (each controller manipulating the *Nova* internal states throughout the REDIS system). This first validation paves the way toward a complete LUC OS leveraging the OpenStack ecosystem.

The remaining of the article is as follows: Section 2 explains our design choices. Section 3 describes the OpenStack and gives first details of our current Proof-of-Concept. Finally Section 5 concludes and discusses future actions.

## 2. THE LUC OS: DESIGN DISCUSSION

The massively distributed cloud we target is an infrastructure that is composed of up to hundreds of micro DCs, which are themselves composed of up to tens of servers. Thus the system in charge of operating such an infrastructure should be able to manage up to thousands of servers spread geographically. Delivering such a system is a tedious task where wrong design choices could prevent to achieve our goal. In this section we first discuss few conceptual considerations that led us to the LUC OS proposal and second remain the major services that the LUC OS should deliver.

### 2.1 From Centralized to Distributed Management

The first way that comes generally to the mind to pilot and use distinct clouds is to rely on classical models like federated approaches: each micro DC hosts and operates its own Cloud infrastructure and a brokering service is in charge of resources provisioning by picking on each cloud. While such approaches can be acceptable for elementary usages, advanced brokering services are mandatory to meet production environment requirements (monitoring, scheduling, automated provisioning, SLAs enforcements ...). In addition to dealing with scalability and single point of failure (SPOF) issues, brokering services should integrate mechanisms similar to those that are already implemented at the level of IaaS managers [4, 15]. Consequently, the development of a brokering solution is as difficult as the one of a IaaS manager but with the complexity of relying only on the least common denominator APIs. While few standards such as OCCI [19] start to be adopted, they do not allow developers to manipulate low-level capabilities of each system, which is generally mandatory to finely administrate resources. In other words, building mechanisms on top of existing ones like in the case of federated systems prevent from going beyond the provided APIs (or require when possible, intrusive mechanisms that must be adapted to the different systems).

The other way to operate such infrastructure is to design and build a dedicated system, *i.e.*, the *LUC Operating System*, in charge of operating all the geographically spread micro DCs in a distributed manner. A LUC OS will define and leverage its own software interface, thus extending capacities of traditional Clouds with its API and a set of dedicated tools. This offers a unique opportunity to go beyond classical federations of Clouds by addressing all crosscutting concerns of a software stack as complex as a IaaS manager and by revising in a fully distributed manner, mechanisms that have been traditionally implemented in a centralized one (service nodes).

The following question is now to analyze whether the collaborations between instances of the system, that is the service nodes, should be structured either in hierarchical way or in a P2P (*i.e.*, flat) one. Few hierarchical solutions have been proposed during the last years in industry [5, 6] and academia [9, 10]. Although they may look easier than P2P structures, hierarchical approaches require additional maintenance costs and complex operations in case of failure. Moreover, mapping and maintaining a relevant tree architecture on top of a network backbone is not meaningful (static partitioning of resources is usually performed). As a consequence, hierarchical approaches do not look to be satisfactory to operate a massively distributed IaaS in-

infrastructure such as the one we target. On the other side, Peer-to-Peer (P2P) file sharing systems are a good example of software that works well at large scale and in a context where computing resources are geographically spread. While largely unexplored for building operating systems, peer-to-peer/decentralized mechanisms have the potential to natively handle the intrinsic distribution of LUC infrastructures as well as the scalability required to manage them. Hence, we propose to leverage advanced P2P mechanisms like overlay networks and distributed hash tables to design the LUC OS building blocks.

## 2.2 Cloud Capabilities

From the administrators and end-users point of views, the LUC OS should deliver a set of high level mechanisms whose assembly results in an operational IaaS system. Recent studies have showed that state of the art IaaS manager [23] were constructed over the same concepts and that a reference architecture for IaaS manager can be defined [20]. This architecture covers primary services that are needed for building the LUC OS :

- The **virtual machines manager** is in charge of managing VMs' cycle of life (configuration, scheduling, deployment, suspend/resume and shut down).
- The **Image manager** is in charge of VM' template files (*a.k.a.* VM images).
- The **Network manager** provides connectivity to the infrastructure: virtual networks for VMs and external access for users.
- The **Storage manager** provides persistent storage facilities to VMs.
- The **Administrative tools** provide user interfaces to operate and use the infrastructure.
- Finally, the **Information manager** monitors data of the infrastructure for the auditing/accounting.

The challenge is thus to propose for each of the aforementioned services, a distributed version of the service that works in a fully decentralized way by relying on advanced P2P mechanisms. However, designing and developing a complete LUC OS from scratch would be an herculean work, including several non-sense actions aiming at simply providing basic mechanisms available in most IaaS solutions. Instead of reinventing the wheel, we propose to minimize both design and implementation efforts by reusing as much as possible existing piece of codes. With this in mind, we propose to investigate whether the OpenStack solution [22] can be revised to fulfill the LUC infrastructure requirements. Concretely, we propose to determine which mechanisms can be directly used and which ones must be revisited with P2P algorithms. This strategy enables us to focus the effort on the key issues such as the distributed functioning, fault tolerance mechanisms, and the organization of efficient collaborations between service nodes of the infrastructure according to the constraints/requirements of the LUC OS.

## 3. REVISING OPENSTACK

OpenStack is an open-source project that aims at developing a complete cloud management system. Similarly to the reference architecture described in the previous Section, it is composed of several services, each one dealing with a particular aspect of a Cloud infrastructure as depicted in Figure 2.




Figure 2: Services composing OpenStack.

OpenStack relies on two kinds of nodes: controller nodes and compute nodes. The formers are in charge of managing and distributing work to the latters that provide computing/storage resources to end-users. In other words, controller nodes correspond to the different services introduced in the previous section while compute nodes are dedicated to the hosting of VMs.

The software architecture of services composing OpenStack is based on the “shared nothing” principles: each controller node (*i.e.*, each service) is connected to the other controllers via two different ways:

- A **messaging queue**, implementing the AMQP protocol that enables the collaboration between sub-services of a controller.
- A **relational database** (DB) that stores inner states of a controller.

Finally, the controllers interact with each others through REST APIs or directly by accessing the inner-states (*i.e.* logical objects used by OpenStack for its functioning) stored in several database backends (mainly MySQL).

Indeed, the process of revising OpenStack, towards a more decentralized and an advanced distributed functioning, can be carried out in two ways: the decentralization of the messaging queue and the use of an alternative to relational databases.

### 3.1 Decentralizing the AMPQ Messaging Queue

OpenStack relies on the RabbitMQ messaging service, which is articulated around the concept of centralized broker. RabbitMQ provide by default a cluster mode, which can be configured to work in a high available mode: several machines, each hosting a RabbitMQ instance, work together in an Active/Active functioning where each queues is mirrored on all nodes. This mode is recommended by the documentation for setting up a multi-controller nodes OpenStack infrastructure. While it has the advantage of being simple, it has the drawback of being very sensible to network latency, and thus is not relevant for multi-site configurations.

According to RabbitMQ documentation, it is possible to use different modes such cluster federation of shovels. This represents a first lead to decentralize the AMQP messaging queue used by OpenStack services. Alternatively, there are few implementations of P2P messaging service such as ActiveMQ [24] or ZeroMQ [14] that would be adapted to the LUC requirements.

### 3.2 Enabling the support of NoSQL databases in Nova

From today's perspective, most of the OpenStack deployments are involving few nodes, thus not requiring more than a single controller node. For larger deployments, involving several controller nodes, there are several possibility considering the architecture of the database: a first possibility would be to use a single database server shared between all the controller nodes, which has the advantage of being a simple solution at the cost of introducing a single point of failure (SPOF) in case the server crashes. The second possibility is to follow recommendations proposed by OpenStack documentation and to deploy on each controller a DB which will be synchronized with the DBs of other servers with instances with a dedicated mechanism [17]. By such a mean, when a controller processes a request and performs some actions on one site, changes in the inner-state are also propagated to all the other locations. From a certain point of view, it gives the illusion that there is only one DB for each service. Although the technique described has been used in different proof-of-concepts, current DB synchronization mechanisms are not scalable enough to cope with a LUC infrastructure deployed on large number of geographical sites.

Another approach would be to replace the DBs used in OpenStack by a more suitable storage backend that would provide a better scalability, like NoSQL databases. Distributed Hash Tables (DHTs) and more recently key/value systems built on top of the DHT concept such as *Dynamo* [8] have demonstrated their efficiency in terms of scalability and fault tolerance properties. On the other hand, there exists NoSQL databases that provide interfaces closer to relational databases, as Cassandra which is fully distributed while supporting queries in a subset of the SQL language (CQL).

#### 3.2.1 Rome library

OpenStack doesn't manipulate directly data stored on the relational database, instead it uses *SQLAlchemy* which serves as an Object-Relational mapping tool by making the link between the relational format used by database and the object format used by programs. Thus OpenStack services manipulates inner-states stored in the database with an object oriented vision.

In light of this, we have developed a python library called "*Rome*" which targets at enabling python programs to manipulate data stored in NoSQL databases with the same programmatic interfaces provided by *SQLAlchemy*. As there exists several types of NoSQL databases (Key/Values stores like Redis, column oriented databases like Cassandra, ...), Rome supports several types of databases, enabling to manipulate them with the same programmatic interfaces.

*Rome* provides the support of relationships in NoSQL databases: relationships between tables are widely used in relational databases thanks to the concept of foreign keys, and most of the NoSQL databases do not implement this. *Rome* support the definition of relationships between model classes, and thus enables to propagate modification on an object to its linked objects.

Moreover, *Rome* implements the *join* technic that enables to combine rows of several tables, in the same manner as *SQLAlchemy* do. While *SQLAlchemy* leverage the "JOIN" operator in SQL language and thus delegate the work to the database server, *Rome* cannot do the same as most of the NoSQL databases (REDIS, Cassandra) do not implement the "JOIN" operation: this operation is executed by "ROME", on the client-side, thus maintaining the compatibility with *SQLAlchemy*.

Finally, *Rome* enables the support of transactions as defined by *SQLAlchemy*: while *SQLAlchemy* rely on transactional mechanisms provided by database server, *Rome* simulates this operation thanks to a distributed lock implemented with *REDIS* that enables to lock model objects and use timeout measurement to detect deadlock situations.

#### 3.2.2 Combining Nova with REDIS

Leveraging the *Rome* library, we have revisited the Nova service which is in charge of managing VMs manager deployed in an OpenStack infrastructure, in order to to replace the current MySQL DB system by *REDIS* [13], a *key/value store*. Technically speaking, we modified the Nova database driver. Indeed, the Nova software architecture has been organised in a way which ensures that each of its sub-services does not directly manipulate the database: they have an indirect access through a service called "nova-conductor" which in turn works with an implementation of the "nova.db.api" programming interface. Developers of Nova provide an implementation of this interface that is using *SQLAlchemy* to manipulate a relational database. We developed a second implementation of this interface that replaces every call to the *SQLAlchemy* by a call to a custom key/value store driver. This enables Nova's services to work with REDIS by only changing the database driver, limiting the level of intrusiveness in the original source code. Thanks to this modification, it is possible to instantiate a distributed cloud and operate it through a single instance of OpenStack composed of several Nova controllers deployed on distinct sites. Figure 4 depicts such a deployment.

Each controller executes a REDIS instance that is configured to work in a clustering way with other instances. One or several controllers can be deployed on each site according to the expected demand in terms of end-users. Finally, a controller can be deployed either on a dedicated node or be mutualized with a compute one as illustrated for Site 3. We highlight that any controller can provision VMs by orchestrating services on the whole infrastructure and not only on the site where it is deployed. Such a behavior is possible thanks to the AMQP bus and the key/value store that go through all controllers.

Finally, it is noteworthy that key/value stores that focus on high-availability and partition tolerance criteria like Cas-

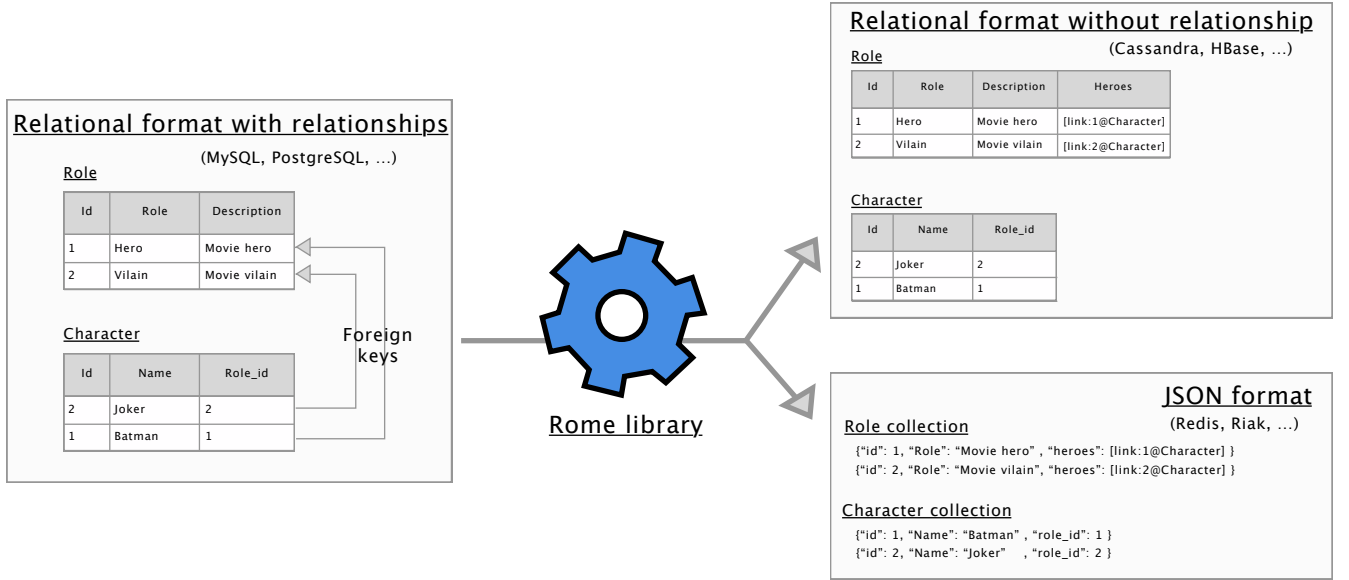


Figure 3: Rome is a library that enables to work with NoSQL databases with an object oriented interface.

sandra [18] would be more appropriate than REDIS for a production deployment. We chose REDIS for its usage simplicity.

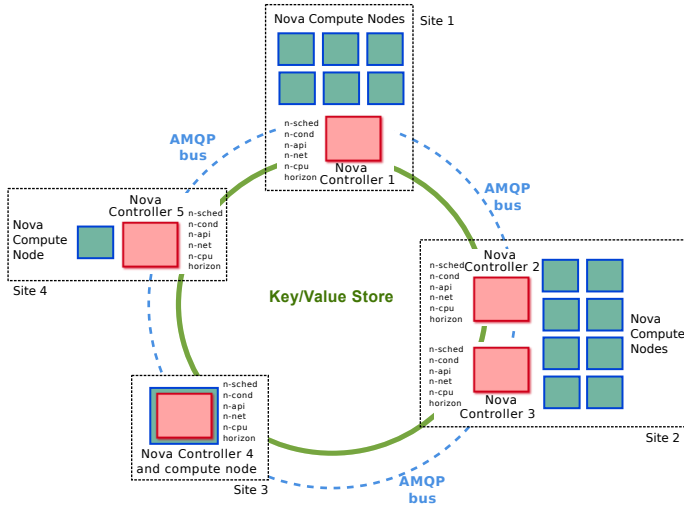


Figure 4: Nova controllers are connected through a shared key/value backend and the AMQP bus.

Our prototype is under evaluation. However, preliminary experiments have been performed throughout 4 sites of Grid'5000 including 12 compute nodes and 4 controllers overall. While this infrastructure was rather small in comparison to our target, it aimed at validating the interconnection of several controllers WANwide and the correct behaviour of OpenStack using our noSQL backend. Our prototype succeeded to provision 500 VMs in 300 seconds (each controller creating 125 VMs in parallel). A second experiment validated the provisioning of 2000 VMs in less than 30 min. We are currently performing comparisons between OpenStack using the historical MYSQL backend *v.s.*, using a key/value store backend. Our goal is to validate that manipulating internal states of Openstack through a noSQL

deliver performances in the same order of the MySQL ones.

#### 4. EXPERIMENTAL VALIDATION ON GRID'5000

The validation of our prototype has been performed thanks to the Grid'5000 testbed [1]. Grid'5000 is a large-scale and versatile experimental testbed for experiment driven research in Computer Science, which enables researchers to get an access to a large amount of computing resources (~ 1000 nodes spread over 10 sites). This delivery of computing resources takes the form of bare-metal machines, on which fully customised software stacks can be deployed, thus giving a very fine control of the experimental conditions.

Various tools have been developed to provide an ease of use, such as monitoring information about networking and power consumption or programming libraries to fine tune each aspect composing an experiment. With this in mind, we developed our prototype using the Execo framework [16] which helped us to deploy and configure each node composing our chosen software stack (Ubuntu 14.04, a modified version of OpenStack "devstack", and the Redis key/value store).

The validation of our work has consisted in two different scenarios. The first scenario consisted in deploying an OpenStack infrastructure composed of 16 compute nodes and a varying number of controllers (between 1 and 4) and to create 500 VM creation requests, fairly distributed among the controllers, in order to detect concurrency problems inherent in using a non relational database backend. The second scenario consisted in creating gradually 2000 VMs on such an infrastructure, in order to detect scalability issues with some tasks involved during the creation of VMs.

During this round of experiments, we compared different OpenStack infrastructures on the basis of several criteria. To enable a fair comparison, each infrastructure was composed of 16 compute nodes (dedicated to the hosting of VMs) and

several controllers and each infrastructure was deployed on top of servers that delivers performance of the same order.

The default database backend that is used for storing states of OpenStack services is MySQL, which is supposed to be deployed on a dedicated node. In the case of this experiment, the MySQL database was hosted on one of the controller nodes.

However using a single node dedicated to the hosting of the database exposes the infrastructure to SPOF problems when the database node crash, a better solution consists in using a replication mechanisms such as Galera, which enables to distribute the MySQL database on several nodes. Its functioning is rather simple: every database node holds the same states, and every write is performed concurrently on every database node. Thus application can request any of the database nodes, the data that is be return will be the same. Thanks to this synchronization mechanism, it is possible to have distributed ACID database that still focus on consistency and availability (CA in the CAP theorem), with the disadvantage of an higher latency due the cost of the synchronization between data stored on nodes.

On the other hand, Redis is a non relational database system which offer native clustering functioning and has opted to prioritise consistency and partition tolerance (CP in the CAP theorem). Redis thus represents an alternative to the use of mechanisms such as MySQL+Galera.

As OpenStack services communicate with each other via RPC calls that are intercepted by an API layer exposed by each instance of a service, we have been able to measure the time take by each call on the API, which give a first overview of the reactivity of a specific OpenStack configuration. Additionnaly, measuring the time taken to spawn a given number of VMs is another way to get an identify the reactivity of a specific configuration.

These two metrics have been measured on several scenarios, each scenario was involving different configurations of the following parameters:

- **Number of controllers**, in order to understand the impact of increasing the number of controllers in an OpenStack deployment. This number has varied between 1 and 4.
- **Database backend**, in order to understand the impact of using a database backend that is different to the relational database used in Nova.
- **Mono-site vs multi-site**, in order to measure how an OpenStack infrastructure was performing when spread accros different geographical sites. In such an experiment, we used 4 sites where each site was hosting 1 controller and 4 computes nodes.

#### 4.1 Creating 500 VMs in parallel

Table 2 presents average of times to serve API requests, while Table 1 introduces the time taken by an OpenStack infrastructure to create 500 VMs in parallel. These two tables presents measures that have been established for several

Database backend	Time (s) taken to create 500 VMs		
	1 CTRL	2 CTRLs	4 CTRLs
MySQL	861.33	497.46	335.21
MySQL +Galera	–	1236.18	966.41
Redis	745.72	533.54	379.67

**Table 1: Time spend to create 500 VMs for monosite based configurations**

Database backend	Time (ms) to serve API requests		
	1 CTRL	2 CTRLs	4 CTRLs
MySQL	35.98	40.95	43.66
MySQL +Galera	–	127.94	175.67
Redis	60.16	79.32	87.65

**Table 2: Average time spend to serve OpenStack API requests for monosite based configurations**

OpenStack configurations by varying two parameters: number of controller nodes involved in the infrastructure and the database backend used to store OpenStack’s states.

According to Table 1, the time taken to create 500 VMs in parallel decrease when the number of controller nodes involved in the OpenStack infrastructure increase, whichever database backend is used to store the OpenStack data states.

A deeper study with the idea of discriminating the different solutions with the performance criterion, will state that our implementation based on the Redis database has performance that are in the same range as those for MySQL: while Redis is 16% faster with one controller, with more controller nodes the gap is closing in favor of MySQL which becomes 13% faster than the Redis solution. This result is explained by the fact that in the 4 controller nodes configuration, the Redis database will be organized in a cluster distributed and synchronized on the 4 controller nodes, while MySQL is running only on one of the controller nodes without any synchronization mechanism. Adding this missing synchronization to MySQL thanks to the Galera software, has a significant impact on the reactivity of the system: with 4 controller nodes a one node MySQL is 289% faster that the MySQL clustered with Galera, which is in turn 255% slower than the Redis based solution.

On the other hand Table 2 shows that increasing the number of controller nodes leads to an increase of the average time take to process an API request. This is due to the fact that increasing the number of controller nodes in an OpenStack configuration leads to increase of the total workload that is required to create the 500 VMs: while the workload necessary to spawn the VMs is distributed among the controller nodes, an extra workload arise partially from the periodic tasks (network, storage, ...) that are run by controller nodes to heal the infrastructure, the remaining extra workload coming from mechanisms to prevent concurrency issues and race condition when manipulating OpenStack states. Furthermore, the addition of controller nodes leads to an increase of the probability that several controller nodes ma-



nipulate a same resource stored in OpenStack's states, thus requiring mechanisms to organize race conditions. In such a situation, a single MySQL node looks better with an average of 43.66 ms to serve API requests in a 4 controller nodes configuration. Clustered database deployments require a significantly higher time to server API requests: in the same 4 controller nodes configuration, Redis is 101% slower than the non clustered MySQL, while the clustered MySQL with Galera is 302% slower than then non clustered one.

In this first set of experiments, a single node MySQL is the more performant solution as it is faster the faster solution in term of time to spawn 500 VMs and in term of time to serve API requests, but at the cost of a lack of synchronization mechanisms which make this solution irrelevant for fault tolerance reasons. If the fault tolerant solutions are considered, using a clustered KVS become a better choice as it doesn't penalize OpenStack functioning and offers better performances than the Galera cluster and provide a fault tolerant functioning.

## 4.2 Creating 2000 VMs gradually

## 5. CONCLUSION AND FUTURE WORK

Distributing the way Cloud are managed is one solution to favor the adoption of the distributed cloud model. In this document, we have presented our view of how such distribution can be achieved. We highlighted that it has however a design cost and it should be developed over mature and efficient solutions. With this objective in mind, we chose to design our system, the LUC Operating System, over OpenStack. This choice presents two advantages: minimizing the development efforts and maximizing the chance of being reused by a large community. As a first step, we modified the Nova SQL backend by a distributed key/value system. Although a more advanced validation of this change is required and the question of which metrics to use remains, this first prototype paves the way toward the distribution of additional OpenStack services.

Among the remaining services, the next candidate is the image service Glance. Indeed, as its images are already stored in fully distributed cloud storage software (SWIFT or CEPH), the next step to reach a fully distributed functioning with Glance is to apply the same strategy that we did with Nova. On the other hand, the situation may be different with some other services: Neutron works with drivers that may not be intended to work in a distributed way. In such situation alternatives will have to be found.

Finally, having a wan-wide infrastructure can be source of networking overheads: some objects manipulated by OpenStack are subject to be manipulated by any service of the deployed controllers, and by extension should be visible to any of the controllers. On the other hand, some objects may benefit from a restrained visibility: if a user has build an OpenStack project (tenant) that is based on few sites, apart from data-replication there is no need for storing objects related to this project on external sites. Restraining the storage of such objects according to visibility rules would enable to save network bandwidth and to settle policies for applications such as privacy and efficient data- replication.

We believe, however, that addressing all these challenges are key elements to promote a new generation of cloud computing more sustainable and efficient. Indeed, by revising OpenStack in order to make it natively cooperative, it would enable Internet Service Providers and other institutions in charge of operating a network backbone to build an extreme-scale LUC infrastructure with a limited additional cost. The interest of important actors such as Orange Labs that has officially announced its support to the initiative is an excellent sign of the importance of our action.

## 6. REFERENCES

- [1] D. Balouek, A. Carpen Amarie, G. Charrier, F. Desprez, E. Jeannot, E. Jeanvoine, A. Lebre, D. Margery, N. Niclausse, L. Nussbaum, O. Richard, C. Pérez, F. Quesnel, C. Rohr, and L. Sarzyniec. Adding virtualization capabilities to the Grid'5000 testbed. In I. Ivanov, M. Sinderen, F. Leymann, and T. Shan, editors, *Cloud Computing and Services Science*, volume 367 of *Communications in Computer and Information Science*, pages 3–20. Springer International Publishing, 2013.
- [2] M. Bertier, F. Desprez, G. Fedak, A. Lebre, A.-C. Orgerie, J. Pastor, F. Quesnel, J. Rouzaud-Cornabas, and C. Tedeschi. Beyond the clouds: How should next generation utility computing infrastructures be designed? In Z. Mahmood, editor, *Cloud Computing, Computer Communications and Networks*, pages 325–345. Springer International Publishing, 2014.
- [3] R. Buyya, R. Ranjan, and R. N. Calheiros. Intercloud: Utility-oriented federation of cloud computing environments for scaling of application services. In *10th Int. Conf. on Algorithms and Architectures for Parallel Processing - Vol. Part I*, ICA3PP'10, pages 13–31, 2010.
- [4] R. Buyya, R. Ranjan, and R. N. Calheiros. Intercloud: Utility-oriented federation of cloud computing environments for scaling of application services. In *Algorithms and architectures for parallel processing*, pages 13–31. Springer, 2010.
- [5] Cascading OpenStack. [https://wiki.openstack.org/wiki/OpenStack\\_cascading\\_solution](https://wiki.openstack.org/wiki/OpenStack_cascading_solution).
- [6] Scaling solutions for OpenStack. <http://docs.openstack.org/openstack-ops/content/scaling.html>.
- [7] CloudStack, Open Source Cloud Computing. <http://cloudstack.apache.org>.
- [8] G. DeCandia, D. Hastorun, M. Jampani, G. Kakulapati, A. Lakshman, A. Pilchin, S. Sivasubramanian, P. Voshall, and W. Vogels. Dynamo: amazon's highly available key-value store. In *ACM SIGOPS Operating Systems Review*, volume 41, pages 205–220. ACM, 2007.
- [9] F. Farahnakian, P. Liljeberg, T. Pahikkala, J. Plosila, and H. Tenhunen. Hierarchical vm management architecture for cloud data centers. In *6th International Conf. on Cloud Computing Technology and Science (CloudCom)*, pages 306–311, Dec 2014.
- [10] E. Feller, L. Rilling, and C. Morin. Snooze: A scalable and autonomic virtual machine management framework for private clouds. In *12th IEEE/ACM Int. Symp. on Cluster, Cloud and Grid Computing (Ccgrid*

- 2012), pages 482–489, 2012.
- [11] J. V. H. Gary Cook. How dirty is your data ? Greenpeace International Report, 2013.
  - [12] A. Greenberg, J. Hamilton, D. A. Maltz, and P. Patel. The cost of a cloud: research problems in data center networks. *ACM SIGCOMM Computer Communication Review*, 39(1):68–73, 2008.
  - [13] J. Han, E. Haihong, G. Le, and J. Du. Survey on nosql database. In *Pervasive computing and applications (ICPCA), 2011 6th international conference on*, pages 363–366. IEEE, 2011.
  - [14] P. Hintjens. *ZeroMQ: Messaging for Many Applications*. ” O’Reilly Media, Inc.”, 2013.
  - [15] I. Houidi, M. Mechtri, W. Louati, and D. Zeghlache. Cloud service delivery across multiple cloud platforms. In *Services Computing (SCC), 2011 IEEE International Conference on*, pages 741–742. IEEE, 2011.
  - [16] M. Imbert, L. Pouilloux, J. Rouzaud-Cornabas, A. Lèbre, and T. Hirofuchi. Using the EXECO toolbox to perform automatic and reproducible cloud experiments. In *1st Int. Workshop on UsiNg and building ClOud Testbeds (UNICO, collocated with IEEE CloudCom*, Dec. 2013.
  - [17] B. Kemme and G. Alonso. Database replication: A tale of research across communities. *Proc. VLDB Endow.*, 3(1-2):5–12, Sept. 2010.
  - [18] A. Lakshman and P. Malik. Cassandra: a decentralized structured storage system. *ACM SIGOPS Operating Systems Review*, 44(2):35–40, 2010.
  - [19] N. Loutas, V. Peristeras, T. Bouras, E. Kamateri, D. Zeginis, and K. Tarabanis. Towards a reference architecture for semantically interoperable clouds. In *Cloud Computing Technology and Science (CloudCom), 2010 IEEE Second International Conference on*, pages 143–150. IEEE, 2010.
  - [20] R. Moreno-Vozmediano, R. S. Montero, and I. M. Llorente. IaaS cloud architecture: From virtualized datacenters to federated cloud infrastructures. *Computer*, 45(12):65–72, 2012.
  - [21] Open Source Data Center Virtualization. <http://www.opennebula.org>.
  - [22] The Open Source, Open Standards Cloud. <http://www.openstack.org>.
  - [23] J. Peng, X. Zhang, Z. Lei, B. Zhang, W. Zhang, and Q. Li. Comparison of several cloud computing platforms. In *2nd Int. Symp. on Information Science and Engineering (ISISE)*, pages 23–27, 2009.
  - [24] B. Snyder, D. Bosnanac, and R. Davies. *ActiveMQ in action*. Manning, 2011.